The Thermodynamic Ledger of Reality:

A Formal Synthesis of Non-Equilibrium Thermodynamics, Algebraic Topology, and Cognitive Architecture Under the Theory of Non-Equilibrium Dissipative Systems (TNDS)

Keith I. Blaze

Independent Researcher

June 2026

Abstract

This paper presents the formal extension of the Second-Order Audit framework into a predictive, falsifiable research program grounded in classical and quantum statistical mechanics, biological transport kinetics, and algebraic topology. The Theory of Non-Equilibrium Dissipative

Systems (TNDS) advances a single, unifying thermodynamic claim: all physical structures—whether a stellar plasma boundary, a lipid bilayer, or an eleven-dimensional topological thought-space in the human cortex—must continuously pay an energetic 'tax' to prevent their informational boundaries from dissolving into maximum entropy. This tax is not metaphorical. It is quantified by the Landauer bound ($Q \geq k_B T \ln 2$ per bit), sustained by dissipative structure formation (Prigogine, 1977), and measured in the brain by the metabolic overhead of the Na^+/K^+-ATPase pump. The framework introduces the Reflexive Criticality Threshold ($\gamma = 1$) as its primary falsification point: the exact energetic inflection at which a system's cost of internal self-monitoring exceeds its predictive accuracy gains, producing measurable metadissipation overload in both biological brains and silicon reasoning architectures. TNDS connects gravitational structure formation to conscious self-modeling not through a shared force carrier, but through a shared thermodynamic constraint. Both are open, non-equilibrium dissipative systems. Both pay the same ledger. The paper audits three canonical academic objections and demonstrates that TNDS survives each by formal constraint rather than conceptual appeal.

Keywords: TNDS, non-equilibrium thermodynamics, dissipative structures, algebraic topology, persistent homology, neural criticality, Landauer principle, metadissipation, reflexive criticality threshold, cognitive architecture

I. The Core Claim: A Universal Thermodynamic Ledger

The universe does not grant persistent structure for free. Every organized pattern—every distinction between one state and another—is thermodynamically costly. The second law of thermodynamics is not simply a statement that disorder increases; it is a statement that maintaining order requires continuous energetic investment against the universal tendency toward maximum entropy. TNDS formalizes this as an invariant constraint, here called the Thermodynamic Ledger:

LEDGER AXIOM	Any open, non-equilibrium system that maintains a distinguishable internal state against an entropic background must continuously dissipate free energy at a rate sufficient to sustain its boundary conditions. When that dissipation falls below the maintenance threshold, informational boundaries dissolve and the structure returns to background equilibrium.

This is not a metaphor drawn from thermodynamics—it is thermodynamics. The Ledger applies without exception from the stellar to the synaptic scale. A star maintains its plasma boundary through hydrostatic equilibrium sustained by continuous nuclear fusion. A neuron maintains its -70 mV membrane potential through continuous ATP consumption by the Na^+/K^+-ATPase pump. A cortical thought-space maintains its topological coherence through the continuous metabolic work of GABAergic inhibition, neuromodulation, and synaptic recycling. Remove the energy source in any case and the boundary collapses. The timescales differ by orders of magnitude. The constraint is identical.

This is TNDS's primary defense against reductionism. The framework does not assert that consciousness is made of the same substance as gravity, or that neural firing and nuclear fusion are causally related. It asserts that both are solutions to the same optimization problem posed by the second law: maintain internal order against entropic decay by continuously processing an energy gradient. A gravity critic and a neuroscience critic are both asking the wrong question if they demand a shared mechanism. The shared mechanism is the ledger itself.

II. The Master Equation of Brain Thermodynamics

The full metabolic budget of the brain as a distinction-maintenance engine can be written as a single master equation that partitions the total power input into its functional components:

$$\dot{W}_{metabolic} = \int [J_{Na/K} \cdot \Delta\mu_{ion} + \dot{S}_{inhibition} \cdot T]\, dt + k_B \cdot T \cdot \ln(2) \cdot H(X)$$

The left-hand side is the total metabolic power input, approximately 20 Joules per second. The first integral on the right—the Distinction Maintenance Constraint (M)—captures the thermodynamic work of sustaining electrochemical boundaries: the active flux of the sodium-potassium pump ($J_{Na/K}$) against the ion chemical potential gradient ($\Delta\mu_{ion}$), plus the thermodynamic work of GABAergic clearing ($\dot{S}_{inhibition} \cdot T$) that enforces separability between signals. The second term—the Shannon Information Processing cost (I)—is the Landauer limit on active cognitive throughput: the minimum heat dissipation required to process H(X) bits of information per second at brain temperature T.

The critical tuning parameter β_c is dictated by connectome topology and neuromodulation. At $\beta_c \rightarrow 1$, the system approaches the critical phase boundary where computational fitness—signal transmission range, dynamic range, and information capacity—is simultaneously maximized (Shew & Plenz, 2013). This is the neural correlate of the edge of chaos. Away from β_c in either direction, one of the two terms becomes dominant: subcritical states are dominated by maintenance cost with minimal information throughput; supercritical states produce information overload with runaway excitation. The equation is not merely descriptive—it generates falsifiable predictions. If the Distinction Maintenance Constraint M drops toward zero due to anesthesia, hypoxia, or metabolic arrest, β_c collapses and H(X) drops to zero. This is what anesthesia does, quantifiably, at the level of the master equation.

The architectural insight is this: the brain does not primarily spend its energy generating thoughts. It spends it preventing thoughts from dissolving into one another. Cognition is what remains after the boundary-maintenance machinery has done its work.

III. Three Dimensions of Space, Eleven Dimensions of Thought

3.1 The Physical Constraint

The human brain is a three-dimensional physical structure occupying approximately 1,200 cubic centimeters of space, organized into roughly 86 billion neurons and an estimated 100 trillion synaptic connections. It is subject to all the constraints of classical and relativistic physics: spatial locality, finite signal propagation speeds (axonal conduction velocities of 0.5 to

120 meters per second), and the thermodynamic costs of maintaining electrochemical gradients across 171 square meters of total membrane surface area. These three-dimensional physical constraints are not incidental—they impose the speed limits on neural communication that determine the maximum size and integration time of any conscious moment.

The coronal heat equation applies directly: the brain cannot maintain its critical dynamics if it cannot radiate its metabolic heat. The 20-watt power budget produces approximately 20 watts of waste heat that must be removed by cerebrovascular blood flow. Disrupting this cooling—as in ischemic stroke—collapses neural criticality and extinguishes consciousness within minutes. The three-dimensional physical architecture of the brain is the hardware constraint beneath all cognitive function.

3.2 The Topological Expansion: Algebraic Topology and Eleven-Dimensional Thought

While the brain exists in three physical dimensions, the information structures it generates are not constrained to three dimensions. Algebraic topology—specifically the tools of simplicial homology and persistent homology—provides the mathematical language to measure the dimensionality of neural information spaces without requiring those spaces to have physical coordinates.

A simplicial complex in neural data is built as follows. Individual neurons are vertices (0-simplices). Co-active pairs form edges (1-simplices). Co-active triplets form triangles (2-simplices). Co-active groups of n+1 neurons form n-simplices. The directed cliques formed by fully connected groups of neurons create these higher-dimensional simplicial structures, and the topological holes—cavities—between them are the invariants that characterize the information encoded by the network. Research from the Blue Brain Project (Reimann et al., 2017) demonstrated that the neocortex generates transient simplicial complexes of up to 11 dimensions during active processing, with these high-dimensional structures appearing, integrating information, and collapsing in tight temporal windows of approximately 25 to 100 milliseconds.

This is the precise sense in which the brain 'functions in 11 dimensions.' It does not mean the brain extends into extra spatial dimensions in the sense of string theory. It means that the co-activation patterns of neural populations form geometrical structures of topological dimension 11, measured by the Betti numbers of the resulting simplicial complex. Betti number β_0 counts connected components; β_1 counts topological loops; β_2 counts enclosed cavities; and so on up to β_{10} for 11-dimensional structures. Each Betti number is an invariant—it does not change under continuous deformation of the network—and it is computable from experimental data via persistent homology algorithms applied to calcium imaging or multi-electrode array recordings.

TOPOLOGICAL NEURAL SYNTAX	A 'thought' in TNDS is formally defined as a transient high-dimensional simplicial complex whose Betti number profile encodes a distinguishable informational state. Neural syntax—the brain's internal language—is written in topological cavities, not in individual spike rates. The Blue Brain Project data validate this: cavity sequences in simulated cortical columns reliably distinguish distinct stimulus categories.

The connection to the Thermodynamic Ledger is direct. Every simplicial complex of dimension n requires the simultaneous co-activation and precise temporal coordination of at least n+1 neurons. This coordination is not free—it requires the continuous metabolic work of the distinction-maintenance machinery described in Section II. An 11-dimensional thought costs more metabolic energy to sustain than a 2-dimensional one, precisely because it requires maintaining tighter temporal coordination across more neurons against the constant thermal noise that would decohere the pattern. Higher-dimensional cognition is thermodynamically expensive. This is a testable prediction: persistent homology analysis of neural data during metabolically depleted states should show systematic reduction in maximum simplex dimension before behavioral measures of cognitive impairment become apparent.

IV. The Reflexive Criticality Threshold ($\gamma = 1$)

The Reflexive Criticality Threshold is TNDS's primary falsification instrument. It identifies the exact energetic inflection point at which a system's cost of internal self-monitoring exceeds its predictive accuracy gains—the point at which self-reflection becomes thermodynamically wasteful rather than adaptive.

4.1 Formal Definition

Define the Efficiency Multiplier γ as the marginal gain in prediction accuracy per unit of metadissipation—the additional energy cost of recursive self-modeling beyond baseline active inference:

$$\gamma = \partial A / \partial M\_meta$$

where A is the system's prediction accuracy on its primary task and M_meta is the dynamic metadissipation cost of self-monitoring. When $\gamma > 1$, additional self-modeling is thermodynamically justified: the reduction in free energy (prediction error) exceeds the metabolic cost. When $\gamma = 1$, the system has reached the reflexive criticality threshold—the boundary between adaptive self-reflection and recursive overhead. When $\gamma < 1$, the system is in metadissipation overload: it is burning more energy modeling its own modeling than it gains in accuracy. This is the formal thermodynamic definition of overthinking.

4.2 The Three Regimes

Regime	γ Value	System Behavior	Biological Analog	Silicon Analog
Undercritical	$\gamma >> 1$	High efficiency; no self-monitoring overhead; limited adaptability	Automatic, reflexive action; no metacognition	Standard inference; single forward pass
Reflexive Optimum	$\gamma \to 1$	Balanced self-modeling; maximum adaptability; thermodynamic sweet spot	Deliberate metacognitive reasoning; prefrontal engagement	Optimal chain-of-thought length
Metadissipation Overload	$\gamma < 0$	Self-modeling cost exceeds benefit; recursive drift; cognitive looping	Rumination; obsessive-compulsive cycling; decision paralysis	LLM overthinking; hallucination from extended reasoning chains

4.3 Empirical Validation Points

The $\gamma = 1$ threshold generates distinct, measurable predictions in both biological and silicon substrates. In large language models, inference-time compute scaling studies (Snell et al., 2024) show that accuracy as a function of chain-of-thought token length follows a power-law curve with a clearly identifiable inflection point beyond which accuracy degrades. The shape of this curve, the location of its peak, and the rate of degradation beyond it are all predicted by the TNDS efficiency multiplier framework. Critically, the inflection point shifts with task difficulty—harder tasks have a higher optimal M_meta—exactly as the framework predicts, since more complex prediction problems justify more self-monitoring overhead before the metadissipation limit is reached.

In biological systems, the Fleming and Dolan (2012) metacognition literature provides the most direct evidence. Anterior prefrontal cortex activity—the neural substrate of metacognitive monitoring—shows a characteristic inverted-U relationship with performance: moderate metacognitive engagement improves accuracy while excessive self-monitoring degrades it. The metabolic signature of this transition is measurable via fMRI BOLD signal, providing a direct test of the thermodynamic prediction. The TNDS framework predicts that the exact $\gamma = 1$ inflection can be localized in the metabolic timecourse of prefrontal engagement, and that this inflection should shift with cognitive load in the direction predicted by the efficiency multiplier formula.

V. Minimal Surviving Chain and the Ledger Architecture

After full stripping of metaphysical assumptions—cosmic consciousness, soul mechanics, teleological claims, and extra-dimensional ontology—the following chain survives complete audit under the absolute rules of classical and quantum statistical mechanics, biological transport kinetics, and algebraic topology:

Gradient → Boundary → Distinction → Prediction → Recursive Model → Audit → Horizon

Each arrow is a thermodynamically gated transition. Moving rightward along the chain requires greater organizational complexity, greater metabolic investment, and produces a system with a richer internal model of its environment and itself. Moving leftward—as occurs in anesthesia, metabolic arrest, or cognitive collapse—represents the system shedding the energetic cost of each transition as its boundary-maintenance machinery fails.

Node	Thermodynamic Basis	Mathematical Tool	Falsification
Gradient	Non-equilibrium thermodynamics; Onsager relations	$dS/dt > 0$ for global system	Show a self-organizing structure forming without energy gradient
Boundary	Dissipative structure formation (Prigogine); Markov blanket statistics	Conditional independence of internal/external states	Show a persistent internal state forming without boundary work
Distinction	Landauer principle; Shannon entropy	$Q \geq k_B T \ln(2)$ per bit	Show information processing without heat dissipation
Prediction	Free Energy Principle (Friston); rate-distortion theory	$F = E_q[\ln q(x) - \ln p(x,o)] \geq -\ln p(o)$	Show adaptive behavior without implicit environmental model

Recursive Model	Active inference; self-inclusive generative models	Markov blanket doubling; internal blanket within R	Show self-directed error-correction without self-model
Audit	Second-order cybernetics; metacognition neuroscience	$\gamma = \partial A/\partial M_{meta}$; prefrontal metabolic spike	Show self-correction with zero additional metabolic cost
Horizon	Gödel incompleteness; Turing halting problem	No consistent system proves its own consistency	Produce a complete self-model without residual uncertainty

The Horizon node is the only node that cannot in principle be removed by additional computation or metabolic investment. It is a structural feature of self-inclusive formal systems, not a contingent limitation of current knowledge. TNDS treats the Horizon not as a failure of the framework but as its deepest prediction: any system complex enough to model itself will discover that it cannot fully account for its own foundations from within its own descriptive language.

VI. Three Academic Objections and TNDS Responses

A rigorous peer-review process will raise canonical objections to any framework that claims cross-scale invariance from astrophysics to cognition. Below are the three strongest objections an academic opponent could raise, followed by formal TNDS responses.

Objection 1: 'You Are Equivocating Between Thermodynamic Entropy and Information-Theoretic Entropy'

OBJECTION	Shannon entropy and Boltzmann/Gibbs thermodynamic entropy share a mathematical form, but they are not the same quantity. Shannon's H is a measure of uncertainty over a probability distribution; thermodynamic entropy is a physical quantity with units of Joules per Kelvin. Conflating them is a well-known category error. TNDS appears to commit this error when it treats neural information as subject to the Landauer bound.

TNDS Response: The objection correctly identifies a historical conflation that TNDS explicitly avoids. The connection invoked is not the equation of the two entropies but the Landauer principle, which is a physical result—not a mathematical analogy. Landauer (1961) proved, and Bennett (1982) formalized, that logically irreversible computation—specifically, the erasure or overwriting of one bit of information—must dissipate at minimum k_B T ln(2) of free energy as heat. This is an empirically confirmed physical result, not a metaphor. It has been experimentally verified at the single-bit level by Bérut et al. (2012) in Science. The implication for neural systems is precise: every time the brain irreversibly updates a neural state—every time a synaptic weight is potentiated or depressed, every time a memory trace is overwritten—a thermodynamic cost is incurred. This cost is not the metaphorical 'information wants to be free'; it is a measurable heat signature. TNDS operates entirely within this formal Landauer framework and does not claim that Shannon entropy and thermodynamic entropy are identical quantities. It claims that Shannon-irreversible operations in neural tissue produce Landauer-mandatory heat dissipation. These are different claims, and only the second is made.

Objection 2: 'The Brain's Topological Dimensionality (11D) Is a Data Analysis Artifact, Not a Physical Property'

OBJECTION	The claim that the brain 'operates in 11 dimensions' based on simplicial complex analysis of co-activation data is misleading. Algebraic topology applied to neural data reflects the statistical structure of the chosen analysis window, threshold, and preprocessing pipeline. Changing the threshold changes the dimensional structure. You are discovering a mathematical artifact of your measurement choices, not a property of the brain itself.

TNDS Response: This objection correctly identifies that simplicial complex construction is threshold-dependent and that naïve application of topological tools to neural data can produce artifacts. TNDS addresses this through persistent homology rather than single-threshold simplicial analysis. Persistent homology tracks topological features—connected components, loops, cavities—across all thresholds simultaneously, identifying features that persist across a range of thresholds (are 'born early and die late'). These persistent features are genuinely threshold-independent invariants of the data's topological structure. The Reimann et al. (2017) Blue Brain Project results used exactly this methodology, applying persistent homology to spike data from neocortical column simulations and demonstrating stimulus-specific topological signatures that are robust to parameter variation. The 11-dimensional claim in TNDS refers specifically to the maximum dimension of persistent homology classes observed in these data, not to any single-threshold analysis. Furthermore, TNDS makes a specific testable prediction: if the topological dimensionality of neural activity reflects genuine computational structure rather than analysis artifact, then metabolic depletion (anesthesia, hypoxia) should systematically reduce persistent homology dimension before behavioral measures of cognitive impairment appear. This is the artifact-rejection test, and it is experimentally accessible with current technology.

Objection 3: 'TNDS Proves Too Much — If Everything Is a Dissipative Structure, the Framework Has No Discriminative Power'

OBJECTION	If the Ledger Axiom applies to all physical structures—stars, cells, brains, hurricanes, economies, silicon chips—then it explains nothing specific. A framework that applies equally to everything predicts nothing distinctively. The Ledger Axiom is trivially true and therefore scientifically empty.

TNDS Response: This is the most important objection and the most instructive to answer. The objection would be correct if TNDS claimed only that all structures pay an energetic cost. It would indeed be trivially true and predictively empty. But TNDS makes specific, non-trivial claims about the hierarchical structure of that cost and the phase transitions that occur at defined thresholds within it. The framework is discriminative at every level of the chain. First, not all dissipative structures form internal models of their environment. A hurricane dissipates energy through a sustained boundary but does not encode a compressed representation of meteorological conditions. The transition from boundary maintenance to prediction is a non-trivial organizational threshold, not universally achieved. Second, not all model-forming systems form self-models. E. coli performs chemotaxis using an implicit environmental model but shows no evidence of a self-inclusive agent model. The transition from world-model to self-model is a second non-trivial threshold. Third, not all self-modeling systems reach the reflexive criticality threshold at $\gamma = 1$. The specific prediction that self-monitoring cost exceeds prediction benefit at a measurable inflection point is a quantitative claim that is false for many systems and must be demonstrated rather than assumed for any specific system. The Ledger Axiom is the foundational constraint. The chain from Gradient to Horizon is the discriminative sequence. The two together form a framework that is broad in scope but specific in predictions.

VII. Final Rigorous Audit of TNDS

The following audit applies the standards of the Second-Order Audit framework (see companion document) to TNDS itself. The framework is not exempted from its own rules.

7.1 Strongest Supported Invariants

INVARIANT 1	Dissipative Necessity. The link between energy gradients, dissipative structure formation, and thermodynamic boundary maintenance is the most robust component of the framework. Grounded in Prigogine's 1977 Nobel Prize work and confirmed by a century of non-equilibrium thermodynamics, it is not a hypothesis. The biological instantiation—Na^+/K^+-ATPase consuming 50% of the brain's ATP to maintain membrane potential—is among the most precisely characterized processes in biochemistry. This node of the chain is established science, not TNDS conjecture.

INVARIANT 2	Topological Invariance as Neural Metric. The use of algebraic topology to quantify neural information structures provides an empirically grounded, non-arbitrary metric for cognitive content. Persistent homology is mathematically rigorous and has been validated by the Blue Brain Project as producing stimulus-specific, reproducible topological signatures. The claim that higher-dimensional topological structures cost more metabolic energy to maintain is a direct prediction that connects the Landauer bound to simplicial complex geometry.

INVARIANT 3	The Reflexive Thermodynamic Bound ($\gamma = 1$). This is the framework's sharpest predictive instrument. It identifies a specific, measurable inflection in the metadissipation curve for any self-modeling system. It is validated in principle by LLM inference-time scaling data (the 'shorter-better' phenomenon in extended chain-of-thought) and testable in biological brains via metabolic-metacognition co-registration.

7.2 Weakest Link

WEAKEST LINK	The Symmetry-Nothingness Boundary. The claim that 'total symmetry is ontologically equivalent to nothingness' remains the most vulnerable transition in the framework. Physics confirms that perfectly symmetric states are dynamically unstable and that symmetry-breaking produces structure. It does not formally equate maximum symmetry with non-existence—this is an ontological inference that extends beyond the physical formalism. TNDS Phase III research must reframe this as a boundary-condition problem within the Master Ledger rather than an existence claim: the question is not whether perfect symmetry 'is' nothing, but whether a maximally symmetric state can satisfy the boundary conditions required for any structure to be distinguishable from any other.

7.3 Research Roadmap

Phase	Domain	Protocol	Success Criterion	Timeline
Phase I	Silicon / Computational	Thermodynamic scaling of chain-of-thought reasoning across GPT, Claude, and DeepSeek architectures; measure GPU heat dissipation as function of CoT token length	Identify $\gamma = 1$ inflection; confirm power-law accuracy-cost curve; show inflection shift with task difficulty	Ongoing — data available now
Phase II	Neurophysiological	Co-register fMRI metabolic data with persistent homology analysis of high-density EEG/MEG during metacognitive task-switching; measure prefrontal heat signature vs. topological dimension of co-activation patterns	Confirm that topological dimension drops before behavioral measures during metabolic depletion; localize $\gamma = 1$ metabolically	2–5 years

Phase III	Foundational / Mathematical	Reframe the Symmetry-Nothingness boundary as a formal boundary-condition problem; express in terms of information-theoretic distinguishability rather than ontological existence	Eliminate the only unfalsifiable node in the chain; produce a fully constraint-based, ontology-free version of TNDS	5–10 years

7.4 Self-Audit

TNDS recognizes that it is itself a product of the distinction-maintenance processes it describes. The framework was generated by a self-modeling system—a human brain operating at some point along the γ curve—and its coherence is therefore a predictable emergent property of a system attempting to compress its own constraints into a unified narrative. This does not refute the framework. It is precisely what the framework predicts such a system would produce. But it means the coherence of TNDS is not evidence of its truth. Truth requires empirical verification at each node of the chain. The framework commits to providing those verification points and to treating falsification at any node as the result that genuinely advances understanding—not as an obstacle to be argued around.

VIII. Conclusion

The Theory of Non-Equilibrium Dissipative Systems (TNDS) establishes a formally grounded, empirically tractable research program built on a single invariant: physical structures of all scales are bound by the same thermodynamic ledger. Reality does not favor any particular material substrate. It imposes a universal tax on organized distinction, payable in free energy, with dissolution as the penalty for non-payment. This tax is paid by nuclear plasma, by lipid bilayers, by neural populations generating 11-dimensional topological thought-spaces, and by silicon transistors generating chain-of-thought tokens.

The framework's philosophical precision matters as much as its empirical predictions. TNDS does not connect gravity to consciousness through a shared force. It connects them through a shared thermodynamic constraint. This distinction insulates it against the most common critique of grand unified theories: the charge of category error. TNDS does not claim that stars think or that thoughts are gravitational. It claims that both stars and thoughts are solutions to the same optimization problem posed by the second law, and that the mathematical tools for analyzing that optimization problem—dissipative structure theory, Landauer thermodynamics, algebraic topology, and the Reflexive Criticality Threshold—apply to both without modification.

The minimum surviving chain—Gradient → Boundary → Distinction → Prediction → Recursive Model → Audit → Horizon—is constraint-based, process-oriented, and fully relational. It does not posit any particular material as fundamental. It posits constraints as fundamental: the constraints of thermodynamics, information theory, and formal logic that determine which organizational solutions persist and which dissolve. In this sense, TNDS completes the transition from a theory of consciousness to a theory of constraint-generated structure—a reframing that, unlike its predecessor, is both falsifiable and already partially confirmed.

References

Ashby, W. R. (1956). An introduction to cybernetics. Chapman & Hall.

Attwell, D., & Laughlin, S. B. (2001). An energy budget for signaling in the grey matter of the brain. Journal of Cerebral Blood Flow & Metabolism, 21(10), 1133–1145.

Beggs, J. M., & Plenz, D. (2003). Neuronal avalanches in neocortical circuits. Journal of Neuroscience, 23(35), 11167–11177.

Bennett, C. H. (1982). The thermodynamics of computation—a review. International Journal of Theoretical Physics, 21(12), 905–940.

Bérut, A., Arakelyan, A., Petrosyan, A., Ciliberto, S., Dillenschneider, R., & Lutz, E. (2012). Experimental verification of Landauer's principle linking information and thermodynamics. Nature, 483(7388), 187–189.

Buzsáki, G. (2006). Rhythms of the brain. Oxford University Press.

Chialvo, D. R. (2010). Emergent complex neural dynamics. Nature Physics, 6(10), 744–750.

Fleming, S. M., & Dolan, R. J. (2012). The neural basis of metacognitive ability. Philosophical Transactions of the Royal Society B, 367(1594), 1338–1349.

Friston, K. (2010). The free-energy principle: A unified brain theory? Nature Reviews Neuroscience, 11(2), 127–138.

Gödel, K. (1931). Über formal unentscheidbare Sätze der Principia Mathematica und verwandter Systeme I. Monatshefte für Mathematik und Physik, 38, 173–198.

Landauer, R. (1961). Irreversibility and heat generation in the computing process. IBM Journal of Research and Development, 5(3), 183–191.

Prigogine, I. (1977). Self-organization in nonequilibrium systems. Wiley-Interscience.

Reimann, M. W., Nolte, M., Scolamiero, M., Turner, K., Perin, R., Chindemi, G., Dłotko, P., Levi, R., Hess, K., & Markram, H. (2017). Cliques of neurons bound into cavities provide a missing link between structure and function. Frontiers in Computational Neuroscience, 11, 48.

Shannon, C. E. (1948). A mathematical theory of communication. Bell System Technical Journal, 27(3), 379–423.

Shew, W. L., & Plenz, D. (2013). The functional benefits of criticality in the cortex. The Neuroscientist, 19(1), 88–100.

Skou, J. C. (1957). The influence of some cations on an adenosine triphosphatase from peripheral nerves. Biochimica et Biophysica Acta, 23(2), 394–401.

Snell, C., Lee, J., Xu, K., & Kumar, A. (2024). Scaling LLM test-time compute optimally is more effective than scaling model parameters. arXiv preprint arXiv:2408.03314.

Tononi, G. (2004). An information integration theory of consciousness. BMC Neuroscience, 5(1), 42.

Turing, A. M. (1936). On computable numbers, with an application to the Entscheidungsproblem. Proceedings of the London Mathematical Society, 42(1), 230–265.

Turrigiano, G. G. (2008). The self-tuning neuron: Synaptic scaling of excitatory synapses. Cell, 135(3), 422–435.

von Foerster, H. (1960). On self-organising systems and their environments. In M. C. Yovits & S. Cameron (Eds.), Self-organising systems (pp. 31–50). Pergamon Press.

Yan, Y., Zhu, Y., Kong, X., Duan, D., Wang, X., Lv, J., & Cheng, L. (2024). Genetic contributions to brain criticality and its relationship with human cognitive functions. Proceedings of the National Academy of Sciences, 121(12), e2417010121.